



IIT KHARAGPUR



NPTEL ONLINE
CERTIFICATION COURSES

Lecture 38: TIMING-DRIVEN ROUTING

PROF. INDRANIL SENGUPTA

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

Timing-Driven Routing

- In modern VLSI chips, interconnects contribute significantly to total signal delay.
 - Interconnect delay is a major concern during routing.
- Timing-driven routing seeks to minimize one or both of:
 - **Maximum sink delay**, which is the maximum interconnect delay from the source node to any sink of a given net.
 - **Total wirelength**, which affects the load-dependent delay of the net's driving gate.


Some Notations

- For a given signal net *net*, let
 - s_0 be the source node
 - $sinks = \{s_1, \dots, s_n\}$ be the sinks
 - $G = (V, E)$ be a weighted graph where $V = \{v_0, v_1, \dots, v_n\}$ represents the source and sink nodes of *net*, and the weight of an edge $e(v_i, v_j) \in E$ represents the routing cost between v_i and v_j .
 - For any spanning tree T over G , let $radius(T)$ be the length of the longest source-sink path in T , and let $cost(T)$ be the total edge weight of T .

An Observation


- Because source-sink wirelength reflects source-sink signal delay, a routing tree ideally minimizes both radius and cost.
- However, for most signal nets, radius and cost cannot be minimized at the same time.
- The next slide illustrates this radius (“*shallow*”) versus cost (“*light*”) tradeoff, where the labels represent edge costs.
 - The first tree has *minimum radius* :: a shortest-path tree that can be constructed using the *Dijkstra’s algorithm*.
 - The second tree has *minimum cost* :: a minimum spanning tree that can be constructed using *Prim’s algorithm*.

- Due to their large cost and large radius, neither of these trees may be desirable in practice.
- The third tree as shown is a *compromise* that has both shallow and light properties.




$radius(T) = 8$
 $cost(T) = 20$

“Shallow”



$radius(T) = 13$
 $cost(T) = 13$

“Light”



$radius(T) = 11$
 $cost(T) = 16$

Tradeoff between
 shallow and light

The Bounded-Radius Bounded-Cost Algorithm

- The BRBC algorithm finds a shallow-light spanning tree with provable bounds on both radius and cost.
 - Each of these parameters is within a constant factor of its optimum value.
- From the graph $G=(V,E)$, the algorithm first constructs a subgraph G' that contains all $v \in V$, and has small cost and small radius.
 - The shortest-path tree T_{BRBC} in G' will also have small radius and cost because it is a subgraph of G' .
 - T_{BRBC} is determined by a parameter $\epsilon \geq 0$, which trades off between radius and cost.
 - When $\epsilon = 0$, T_{BRBC} has *minimum radius*; when $\epsilon = \infty$, T_{BRBC} has *minimum cost*.

- T_{BRBC} satisfies both the following inequalities:

$$radius(T_{BRBC}) \leq (1 + \varepsilon) \cdot radius(T_S)$$

$$cost(T_{BRBC}) \leq \left(1 + \frac{2}{\varepsilon}\right) \cdot cost(T_M)$$

where T_S is a shortest-path tree of G , and T_M is a minimum spanning tree of G .

Prim-Dijkstra Tradeoff

- This method trades off radius and cost by using an explicit, quantitative metric.
- The minimum cost and minimum radius objectives are typically optimized using Prim's algorithm and Dijkstra's algorithm.
 - Though they target different objectives, these algorithms construct spanning trees over the set of terminals in very similar ways.
 - From the set of sinks S , each algorithm begins with tree T consisting only of s_0 , and iteratively adds a sink s_j and the edge connecting a sink s_i in T to s_j .
 - The algorithms differ only in the cost function by which the next sink and edge are chosen.

- In Prim's algorithm, sink s_j and edge $e(s_i, s_j)$ are selected to minimize the edge cost between sinks s_i and s_j , where $s_i \in T$ and $s_j \in S - T$.


$$\text{cost}(s_i, s_j)$$
- In Dijkstra's algorithm, sink s_j and edge $e(s_i, s_j)$ are selected to minimize the path cost between source s_0 and sink s_j , where $s_i \in T$ and $s_j \in S - T$.

$$\text{cost}(s_0, s_i) + \text{cost}(s_i, s_j)$$
- To combine the two objectives, **PD Tradeoff** iteratively adds the sink s_j and the edge $e(s_i, s_j)$ to T such that

$$\gamma \cdot \text{cost}(s_0, s_i) + \text{cost}(s_i, s_j)$$
is minimum over all $s_i \in T$ and $s_j \in S - T$ for a prescribed constant $0 \leq \gamma \leq 1$.

The Tradeoff


- When $\gamma = 0$, the PD tradeoff is identical to Prim's algorithm, and T is a minimum spanning tree T_M .
- As γ increases, the PD tradeoff constructs spanning trees with progressively higher cost but lower radius.
- When $\gamma = 1$, the PD tradeoff is identical to Dijkstra's algorithm, and T is the shortest-paths tree T_S .



$$\text{radius}(T) = 19$$

$$\text{cost}(T) = 35$$

$$\gamma = 0.25$$



$$\text{radius}(T) = 15$$

$$\text{cost}(T) = 39$$

$$\gamma = 0.75$$

Minimization of Source to Sink Delay

- For a multi-pin net, the sink with the least timing slack is called the *critical sink* of the net.
 - A routing tree construction that does not consider critical-sink information may create avoidable negative slack, and degrade the overall timing performance of the design.
 - Several routing tree construction approaches exist that address the critical-sink routing problem.
- We iteratively form a tree by adding sinks, and optimizes the critical sinks.

Critical-sink Routing Tree (CSRT) Problem

- Given a signal net with source s_0 , sinks $S = \{s_1, s_2, \dots, s_n\}$, and sink criticalities $\alpha(i) \geq 0$, for each $s_i \in S$ construct a routing tree T such that

$$\sum_{i=1}^n \alpha(i) \cdot t(s_0, s_i)$$

is minimized, where $t(s_0, s_i)$ is the signal delay from source s_0 to sink s_i , and the sink criticality $\alpha(i)$ reflects the timing criticality of the sink s_i .

- If a sink is on the critical path, then its timing criticality will be greater than that of other sinks.

Critical-sink Steiner Tree Heuristic

- This heuristic first constructs a heuristic minimum-cost Steiner tree T_0 over all terminals of S except the critical sink s_c .
- Then, to reduce $t(s_0, s_c)$, the heuristic adds s_c into T_0 by heuristic variants:
 - H_0 : introduce a single wire from s_c to s_0 .
 - H_1 : introduce the shortest possible wire that can join s_c to T_0 , so long as the path from s_0 to s_c is monotone (i.e. shortest possible total length).
 - H_{best} : try all shortest connections from s_c to edges in T_0 as well as from s_c to s_0 . Perform timing analysis on each of these trees and return the one with the lowest delay at s_c .

RAT Tree Problem

- For a signal net with source s_0 and sink set S , find a minimum-cost routing tree T such that

$$\min (RAT(s) - t(s_0, s)) \geq 0, \text{ for all } s \in S$$

where $RAT(s)$ is the required arrival time for sink s , and $t(s_0, s)$ is the signal delay in T from source s_0 to sink s .

END OF LECTURE 38

NPTEL





IIT KHARAGPUR



NPTEL ONLINE
CERTIFICATION COURSES

Lecture 39: PHYSICAL SYNTHESIS (PART 1)

PROF. INDRANIL SENGUPTA

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

Physical Synthesis

- Physical synthesis is a collection of timing optimizations to fix negative slack.
 - Consists of creating timing budgets and performing timing corrections.
- Timing budgets include:
 - Allocating target delays along paths or nets.
 - Often done during placement and routing stages.
 - Can also be done during timing correction operations.

- Timing corrections can be done using various techniques:
 - a) Gate sizing
 - b) Buffering
 - c) Netlist restructuring

(a) Gate Sizing


- Modern-day VLSI chip design is predominantly based on the standard-cell based design style.
 - Each logic gate (e.g. NAND or NOR) is typically available in multiple sizes that correspond to different drive strengths.
 - The drive strength is the amount of current that the gate can provide during switching.
- When a gate drives other gates, the gate output node faces some load capacitance depending on the fanout and sizes of the driven gates.
 - Illustrative example with CMOS gates.

- Suppose that a gate v has three versions A , B and C of differing sizes:
 $size(v_A) < size(v_B) < size(v_C)$
- A gate with larger size will have lower output resistance and can drive a larger load capacitance with smaller load-dependent delay.
- However, a larger gate also has larger intrinsic delay due to the parasitic output capacitance of the gate itself.
 - When the load capacitance is large, the load-dependent delay will dominate:
 $t(v_C) < t(v_B) < t(v_A)$

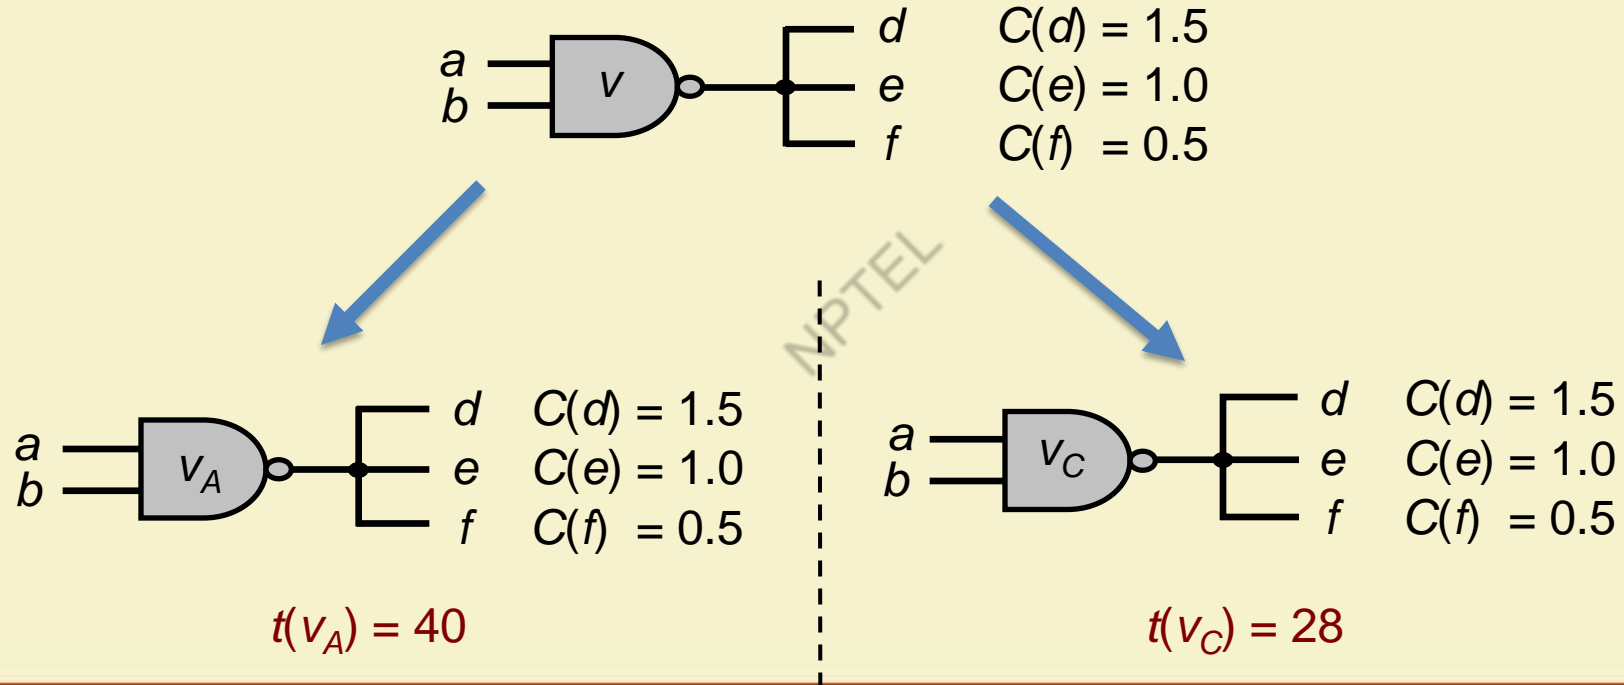
- When the load capacitance is small, the intrinsic delay will dominate:

$$t(v_A) < t(v_B) < t(v_C)$$

- Increasing *size(v)* also increases the gate capacitance of *v*, which in turn increases the load capacitance seen by the fanin drivers.




- Resizing transformations adjust the size of a gate v to achieve a lower delay.



- The total load capacitance driven by gate v is $C(d)+C(e)+C(f) = 3 \text{ fF}$.
- Gate delays are calculated based on the load-delay relations shown in the graph.
- For a load capacitance of 3 fF, gate delay reduces from 40 ps to 28 ps if v_C is used instead of v_A .



(b) Buffering


- A buffer is a gate, typically two serially-connected inverters, that regenerates a signal without changing the functionality.



- Buffers can improve delays by:
 - Speeding up the circuit or by serving as delay elements.
 - Modify transition times to improve signal integrity and coupling-induced delay variation.
 - Shielding capacitive load on the output side.

- Drawbacks of using buffers:
 - They consume more area.
 - They increase power consumption.
- In spite of judicious use of buffering by modern EDA tools, the number of buffers has been steadily increasing in large designs due to technology scaling trends.
 - Interconnect delays dominate as compared to gate delays.
- In modern high-performance designs, buffers can comprise 10-20% of all standard cell instances, and more than 40% in some designs.





- In the example, the (actual) arrival time at fanout pins d-h for gate v_B is $t(v_B) = 45\text{ ps}$.
 - Let pins d and e be on the critical path with RAT's below 35 ps .
 - Let the input pin capacitance of buffer y be 1 fF .
- Adding y reduces the load capacitance of v_B from 5 to 3 , and reduces the arrival times at d and e to $t(v_B) = 33\text{ ps}$.
- After y is inserted, the arrival times at pins f , g and h becomes $t(v_B) + t(y) = 33 + 33 = 66\text{ ps}$.

END OF LECTURE 39



IIT KHARAGPUR



NPTEL ONLINE
CERTIFICATION COURSES

Lecture 40: PHYSICAL SYNTHESIS (PART 2)

PROF. INDRANIL SENGUPTA

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

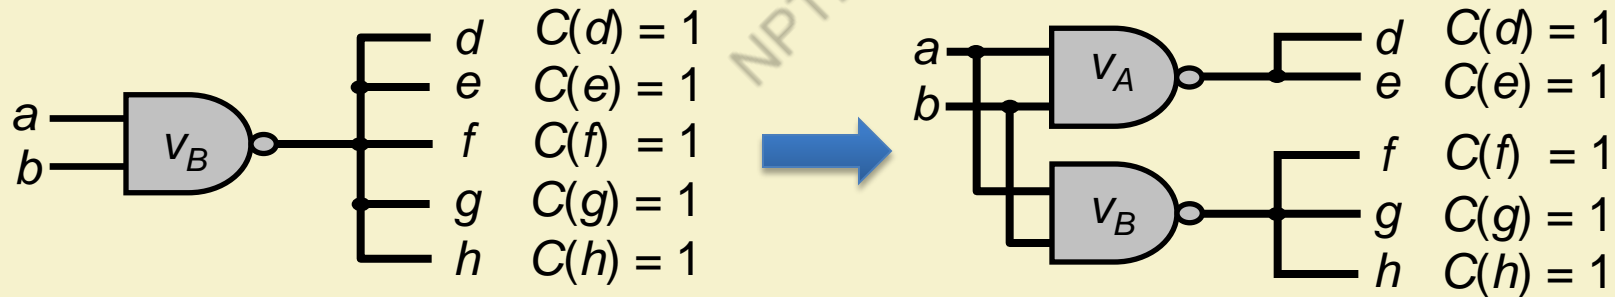
(c) Netlist Restructuring


- Here we modify the netlist itself to improve timing.
 - Such changes must not alter the circuit functionality, but can use additional gates or modify (rewire) the connections between existing gates to improve driving strength and signal integrity.
- Some common netlist modification techniques include:
 - *Cloning*: duplicate gates
 - *Redesign fanin or fanout tree*: changing the topology of gates
 - *Swapping communicative pins*: changing the connections
 - *Gate decomposition*: e.g. changing AND-OR to NAND-NAND
 - *Boolean restructuring*: applying Boolean algebra rules to change circuit gates

Cloning

- Gate duplication can reduce delay in two situations:
 - When a gate with significant fanout may be slow due to its fanout capacitance.
 - When a gate's output fans out in two different directions, making it impossible to find a good placement for this gate.
- The effect of cloning is to split the driven capacitance between two equivalent gates, at the cost of increasing the fanout of upstream gates.
- The second application of cloning allows the designers to replicate gates and place each clone closer to its downstream logic.


- Using the same load-delay relation as shown earlier, the initial gate delay $t(v_B) = 45 \text{ ps}$.
- After cloning, $t(v_A) = 30 \text{ ps}$ and $t(v_B) = 33 \text{ ps}$.





- When the downstream capacitance is large, buffering may be a better alternative than cloning because buffers do not increase the fanout capacitance of upstream gates.
- However, buffering cannot replace placement-driven cloning.


- The gate v drives five signals $d-h$, where signals $d-f$ are close, and g & h are located much farther away.
- To mitigate the large interconnect delay, the gate is cloned and a new copy v' is placed closer to g & h .



Redesign of Fanin Tree



- Logic design often provides a netlist with the minimum number of logic levels.
- Minimizing the maximum number of gates on a path between sequential elements tends to produce a balanced circuit with similar path delays from inputs to outputs.
 - However, input signals may arrive at varied times, so the level minimized circuit may not be timing-optimal.

- The arrival time $AAT(f) = 6$ in the original circuit.
- The modified unbalanced network has a shorter input-output path for the latest arriving signal.



Redesign of Fanout Tree


- It is possible to improve timing by rebalancing the output load capacitance in a fanout tree so as to reduce the delay of the longest path.
- In the example shown on the next slide, the buffer y_1 is needed because the load capacitance of critical path $path_1$ is large.
- By redesigning the fanout tree to reduce the load capacitance of $path_1$, use of buffer y_1 can be avoided.
 - Increased delay on $path_2$ may be acceptable if that path is not critical even after the load capacitance of buffer y_2 is increased.



Swapping Commutative Pins


- Although the input pins of some gates (e.g. a NAND gate) are logically equivalent, in the actual transistor-level netlist they can have different delays to the output pin.
 - Path delays can change when the input pin assignment is changed.
- Rule-of-thumb:
 - Assign a later (sooner) arriving signal to an equivalent input pin with shorter (longer) input-output delay.

- The internal timing arcs are labeled with corresponding delays in parentheses, and the pins are labeled with corresponding arrival times in angular brackets.
 - The arrival time at f can be improved from 5 to 3 by swapping pins a and c .




Gate Decomposition

- In CMOS, a gate with multiple inputs usually has larger size and capacitance, as well as a more complex transistor-level network.
- Decomposition of multiple-input gates into smaller, more efficient gates can decrease delay and capacitance.
- An example is shown on the next slide, where two equivalent NAND-level networks are shown.




Boolean Restructuring

- Rules of Boolean algebra can be used to implement a function in various ways.
- An example is shown on the next slide.
 - The distributive law $f(a,b,c) = (a + b)(a + c) = a + bc$ is exploited to improve timing.
 - The two functions $x = a + bc$ and $y = ab + c$ are realized with signal arrival times $AAT(a) = 4$, $AAT(b) = 1$, and $AAT(c) = 2$.
 - When implemented with a common node $a + c$, we get $AAT(x) = AAT(y) = 6$.
 - Implementing x and y separately achieves $AAT(x) = 5$ and $AAT(y) = 6$.



$$x(a,b,c) = (a + b)(a + c)$$

$$y(a,b,c) = (a + c)(b + c)$$



$$x(a,b,c) = a + bc$$

$$y(a,b,c) = ab + c$$

END OF LECTURE 40

NPTEL





IIT KHARAGPUR



NPTEL ONLINE
CERTIFICATION COURSES

Lecture 41: PERFORMANCE-DRIVEN DESIGN FLOW

PROF. INDRANIL SENGUPTA

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING




Introduction

- We have discussed several techniques to improve timing of digital circuits.
- All these optimization techniques can be combined in a consistent *performance-driven physical design flow*.
- Due to the nature of performance optimizations, their ordering is important.
 - Evaluation steps (like STA) must be invoked several times.
 - Some optimizations like buffering must be redone multiple times to facilitate a more accurate evaluation.

Typical Physical Design Flow

- A typical design flow starts with *chip planning*.
 - Includes *I/O placement*, *floorplanning*, and *power planning*.
- *Trial synthesis* provides the floorplanner with an estimate of the total area needed by modules.
 - Besides logic area, additional space must be provided to account for buffers, routing, and gate sizing.
- Then *logic synthesis* and *technology mapping* produce a gate-level or cell-level netlist from a high-level specification, tailored to a specific technology library.

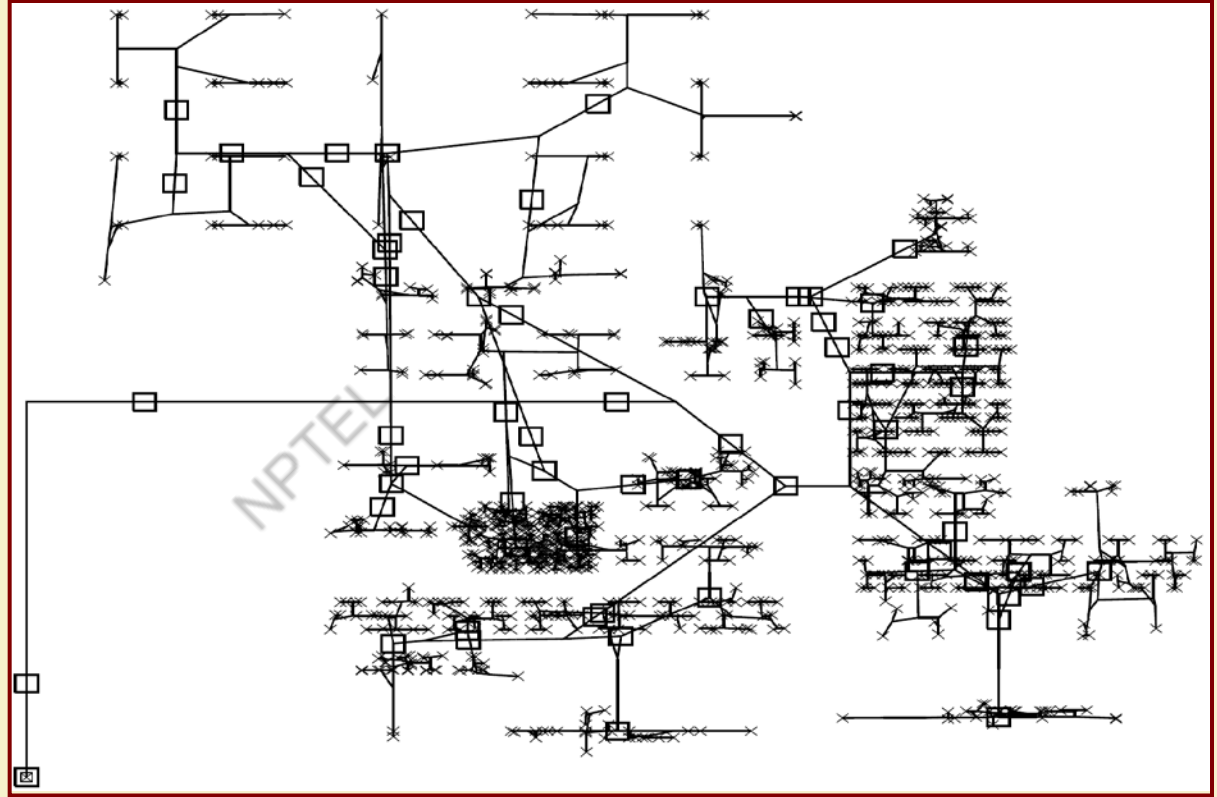
- Next, *global placement* assigns locations to each movable object.
 - Most of the cells are clustered in highly concentrated regions (colored black).
 - As the iteration proceeds, the cells are gradually spread across the chip, such that they no longer overlap (colored light gray).



- After global placement, the sequential elements are finalized.
 - Then a *clock network* is generated, either in the form of a clock tree, or in the form of a network consisting of trees and meshes.





NPTEL

Buffered Clock Tree in a Small CPU Design

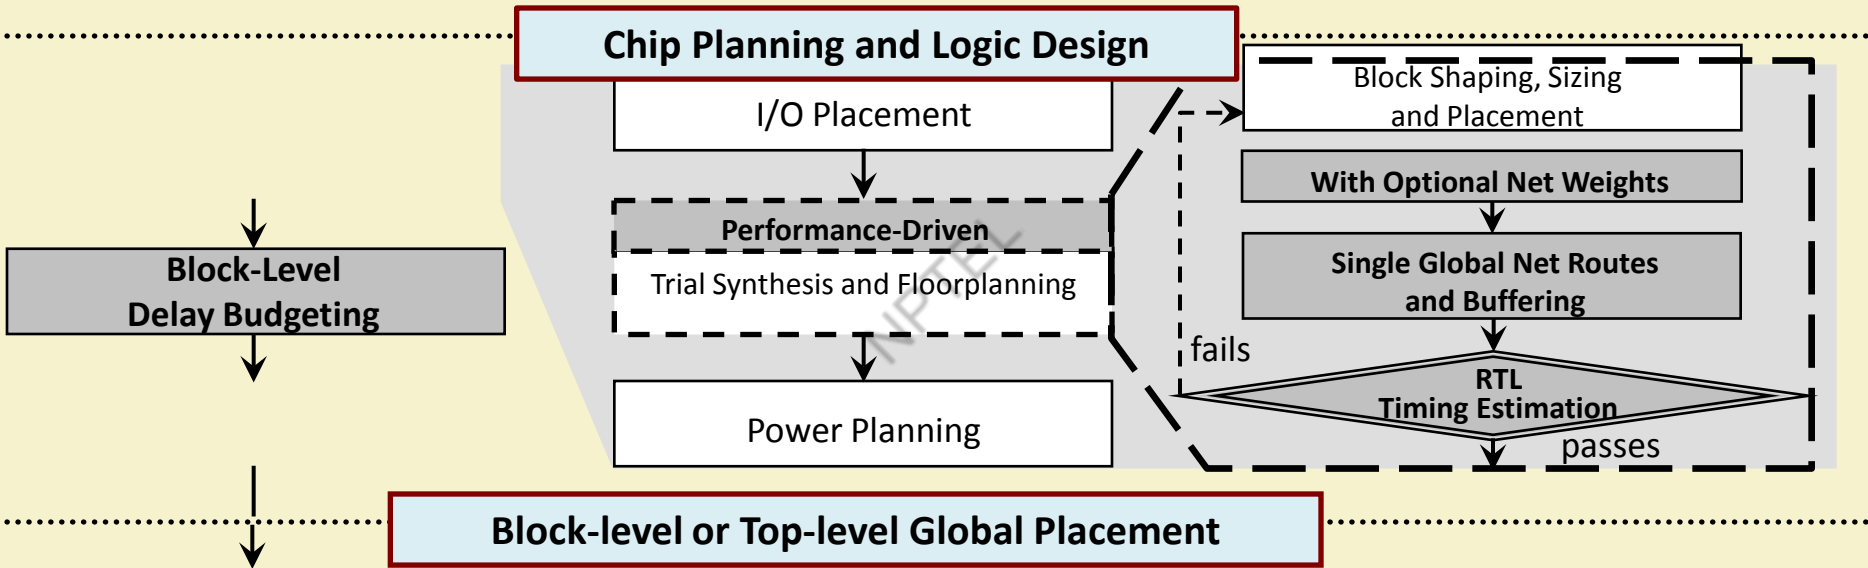


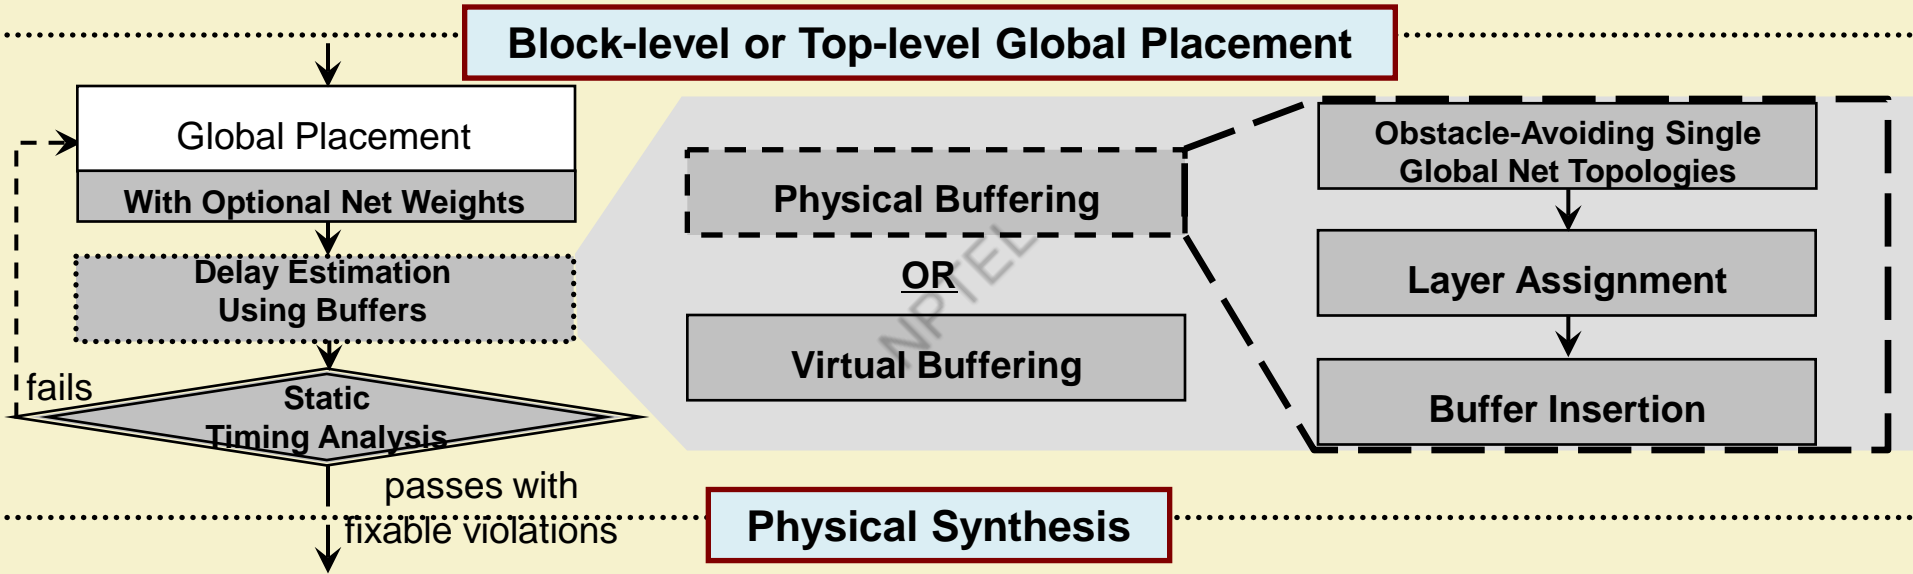
- The locations of globally placed cells are first temporarily rounded to a uniform grid, and then these rounded locations are connected during *global routing* and *layer assignment* (where each route is assigned to a specific metal layer).
- The routes indicate areas of wiring congestion.
 - This information is used to guide *congestion-driven detailed placement*.

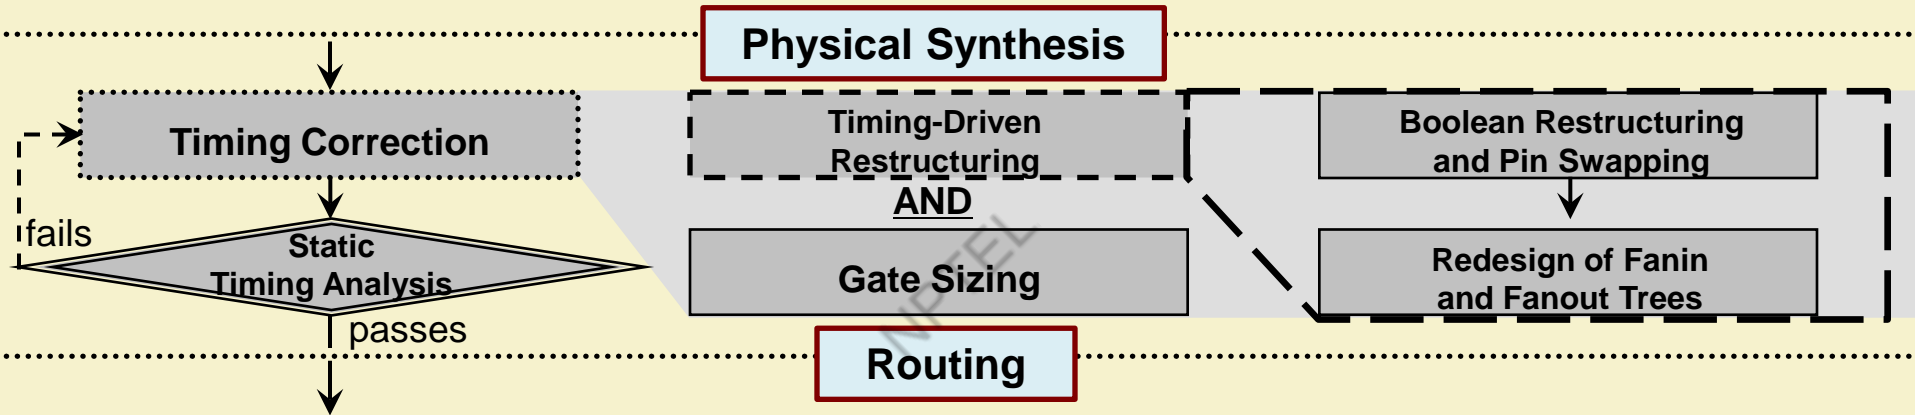
Progression of Congestion Maps Through Iterations of Global Routing

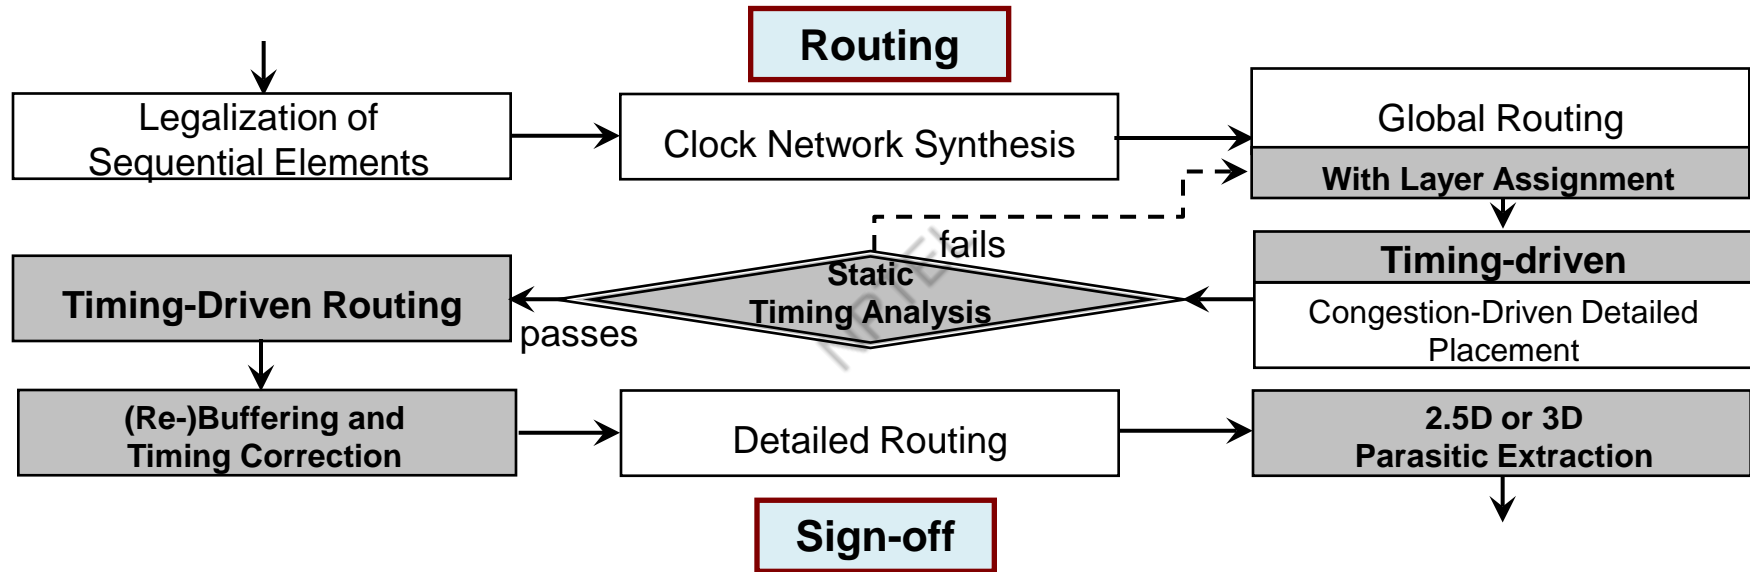


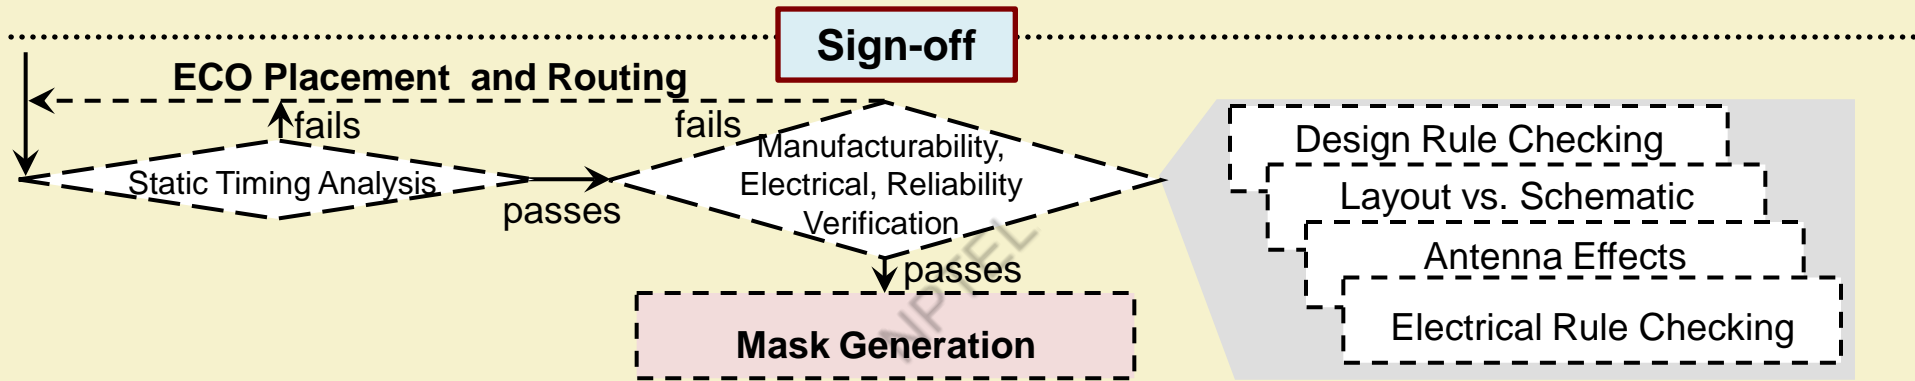
- The global routes of signal nets are then assigned to physical routing tracks during *detailed routing*.
- The layout generated during the place-and-route stage is subjected to *reliability, manufacturability, and electrical verification*.
- During *mask generation*, each standard cell and each route are represented by collections of rectangles in a format suitable for generating optical lithography masks for chip fabrication.











Engineering Change Order (ECO):: how last minute changes are incorporated in a design.

END OF LECTURE 41

NPTEL





IIT KHARAGPUR



NPTEL ONLINE
CERTIFICATION COURSES

Lecture 42: MISCELLANEOUS APPROACHES TO TIMING OPTIMIZATION

PROF. INDRANIL SENGUPTA

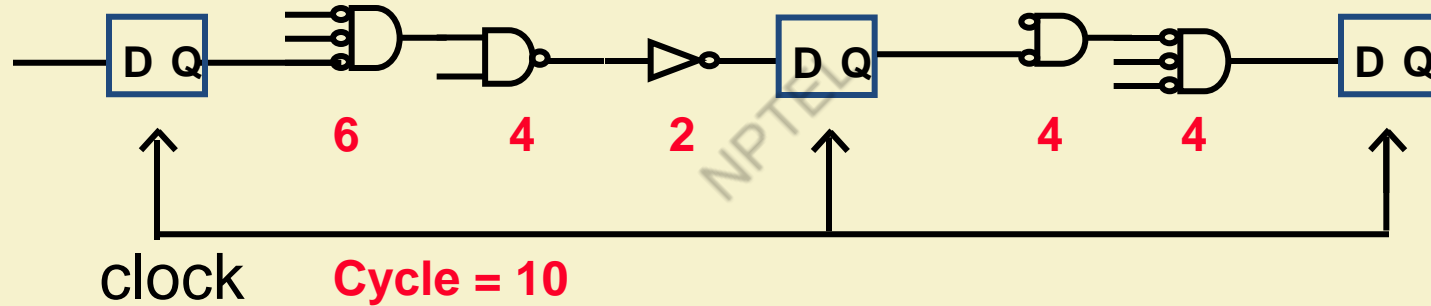
DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

Post-Route Optimization Approaches

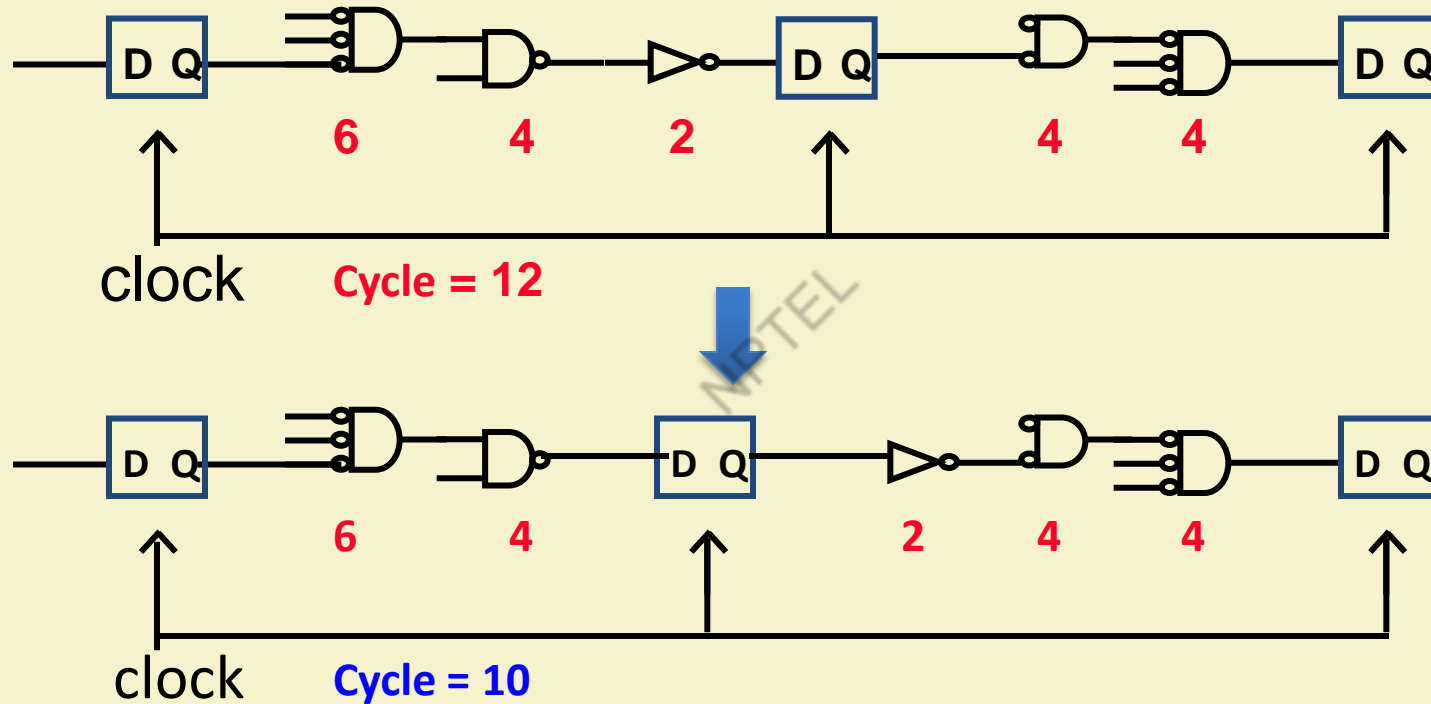
- *Retiming* and *Useful Clock Skew* based approaches.
 - Make modifications to the circuit netlist so as to meet timing requirements.
 - To be explained with the help of illustrative examples.

Re-Timing

- How to meet the 10 ns clock cycle time requirement?

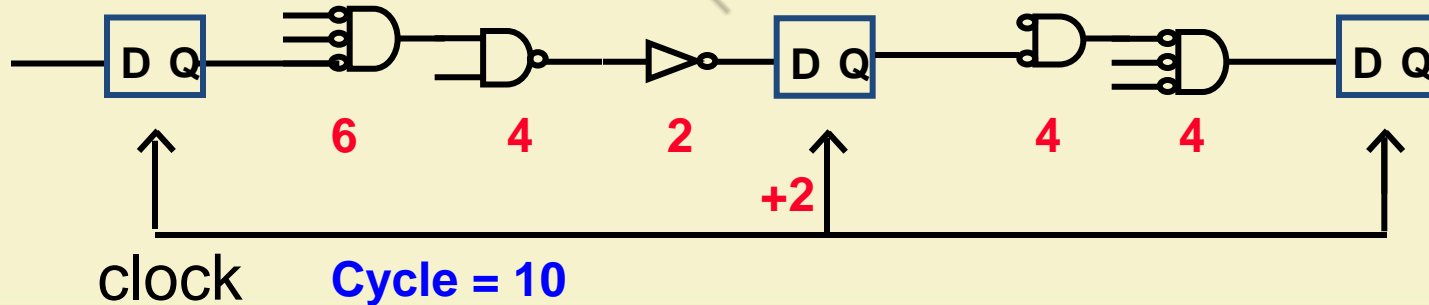


- Re-order sequential elements and combinational logic.

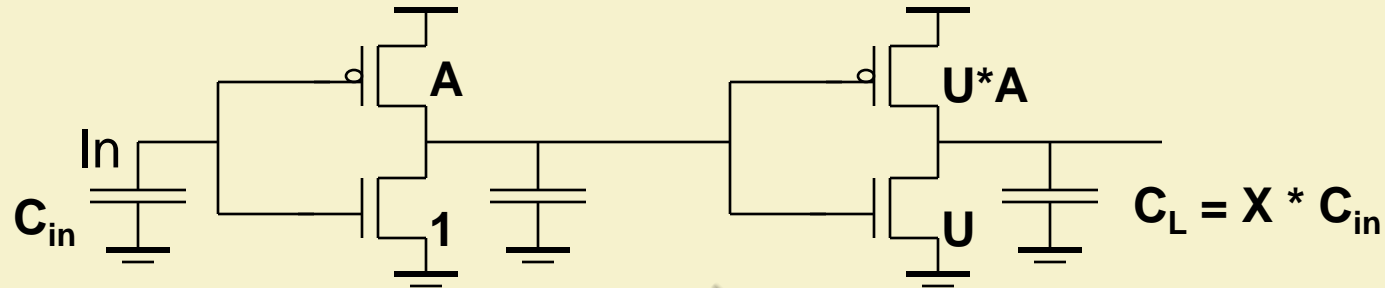


Useful Clock Skew

- Adding skew to a clock signal to meet clock period requirement.
- Several ways:
 - Insert delay cells
 - Add dummy capacitive load
 - Snaking of interconnects




Driving Large Capacitances: Inverter as Buffer



- Total propagation delay = $t_p(\text{inv}) + t_p(\text{buffer})$
- t_{p0} = delay of min-size inverter with single min-size inverter as fanout load
- Minimize $t_p = U \cdot t_{p0} + X/U \cdot t_{p0}$
 $U_{\text{opt}} = \sqrt{X} ; t_{p,\text{opt}} = 2 t_{p0} \cdot \sqrt{X}$
- Use only if combined delay is less than that for unbuffered case.


Delay Reduction with Cascaded Buffer



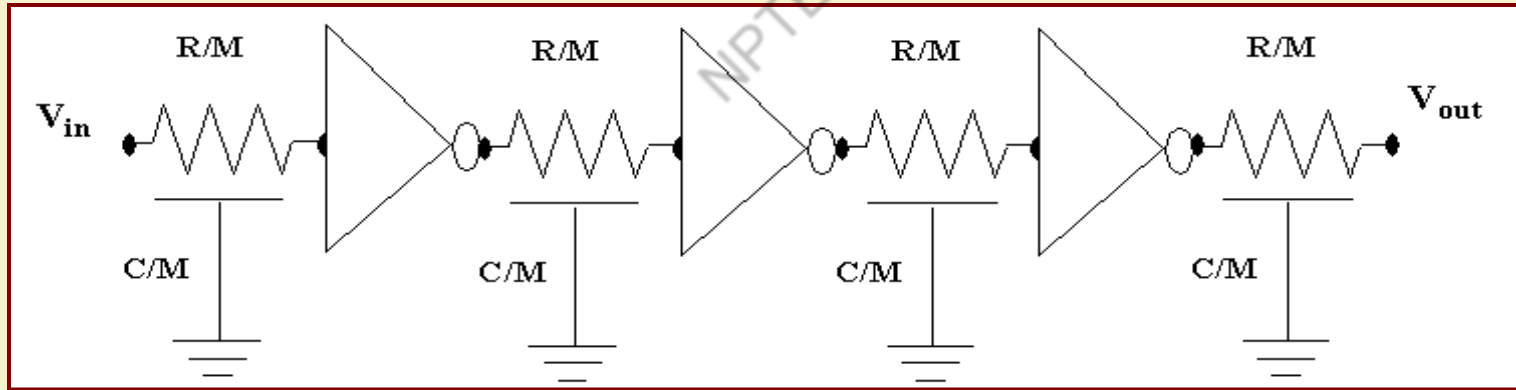
- Cascaded buffers with increasing sizes can reduce delay (*u: scaling factor*).
- If load is driven by a large transistor (which is driven by a smaller transistor), then its turn-on time dominates overall delay.
- Each buffer charges the input capacitance of the next buffer in the chain and speeds up charging, reducing total delay.
- Cascaded buffers are useful when $R_{int} < R_{tr}$.

Reducing RC Delay with Repeaters

- RC delay is quadratic in length \rightarrow must reduce length of interconnects.
- Observation: $2^2 = 4$ and $1 + 1 = 2$, but $1^2 + 1^2 = 2$.




- *Repeater* :: strong driver (usually inverter of pair of inverters) that is placed along a long RC line to *break-up* the line and reduce delay.



Repeaters versus Cascaded Buffers

- Repeaters are used to drive long RC lines.
 - Breaking up the quadratic dependence of delay on line length is the goal.
 - Typically sized identically.
- Cascaded buffers are used to drive large capacitive loads, where there is no parasitic resistance.
 - We put all buffers at the beginning of the load.
 - Useless for a long RC wire since the wire RC delay would be unaffected and would dominate the total delay.

- Optimum buffering for a uniform long interconnect.
 - Cascaded buffers at source and sink.
 - Identical sized and spaced repeaters in between.



END OF LECTURE 42

NPTEL

